

date-a-scientist_Carlos_solution

May 20, 2021

1 Portfolio Project: OKCupid Date-A-Scientist

1.1 Project Author: Carlos Paiva

1.2 Introduction

The objective of this project is to analyze the data from OKCupid, an online dating application, and find patterns inside the data provided on the platform by its users similar to the exercise OKCupid does to identify good matches amongst their members. Like OKCupid, many other apps use sophisticated data science techniques to recommend possible matches and to optimize the user experience, and this project will leverage on the vast amount of information that they have access to.

1.3 Scoping

This project will be divided in four sections. First, the overall project goals will be defined as well as the intention of the project. Secondly follows the explanation of the source of the data including a basic description of the information contained within it. Then, the project will describe the analysis, methods and tools used to achieve the project goals. Finally, some conclusions and recommendations will be drawn from the analysis of the obtained results.

1.3.1 Project Goals

This project has 2 main goals. The first goal is to provide some problem solving, data analysis, and coding skills to the aspiring Data Scientist in the context of the respective Career Path in Codecademy. Second, to answer a critical research question: Can OKCupid's use some variables from the users profile to accurately predict their zodiac sign when this information is not available? Knowing the zodiac sign of each individual is key since many users think that their astrological sign has an important role when choosing a romantic partner.

Therefore, this project defines the problem to solve as the request to build a machine learning model for OKCupid to predict the missing zodiac signs as accurately as possible, by selecting the dataset variables of more relevance to increase the model accuracy.

1.3.2 Data

The dataset used for the project analysis (`profiles.csv`) has been provided by Codecademy.com. Before having any modifications from the project side, it contains 59946 rows and 31 columns in csv format, which will be transformed into a dataframe. This dataset contains information about the age, body type, variety of diet, drinking frequency, drugs-consumption frequency, education

level, topic of essays 0 to 9, ethnicity, height, income, job, last time online, location, whether they have children or not, sexual orientation, pets, religion, sex, zodiac sign, whether smoker or not, language and current status.

1.3.3 Analysis

This project will leverage on descriptive statistics and visualization tools to fully understand the data before building 4 different machine learning classification algorithms (K-Nearest Neighbors, Decision Tree, Random Forest, and Naive Bayes) to obtain zodiac sign predictions based of the data from the remaining columns.

1.3.4 Evaluation

Finally, this project will evaluate the accuracy of the 4 different ML models to select the one with the highest accuracy. Recommendations for improving the study, as well as next steps, will also be mentioned.

1.4 Data

1.4.1 Project Setup

```
[52]: # Importing basic Python libraries:
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns

# Reading csv files as pandas dataframe:
profiles = pd.read_csv('profiles.csv')

# Printing original dataset first 5 rows:
profiles.head()
```

```
[52]:
```

	age	body_type	diet	drinks	drugs	\
0	22	a little extra	strictly anything	socially	never	
1	35	average	mostly other	often	sometimes	
2	38	thin	anything	socially	NaN	
3	23	thin	vegetarian	socially	NaN	
4	29	athletic	NaN	socially	never	

	education	\
0	working on college/university	
1	working on space camp	
2	graduated from masters program	
3	working on college/university	
4	graduated from college/university	

	essay0	\
0	about me: \n \ni would love to think...	

```

1 i am a chef: this is what that means.<br />\n1...
2 i'm not ashamed of much, but writing public te...
3     i work in a library and go to school. . .
4 hey how's it going? currently vague on the pro...

                                essay1 \
0 currently working as an international agent fo...
1 dedicating everyday to being an unbelievable b...
2 i make nerdy software for musicians, artists, ...
3     reading things written by old dead people
4         work work work work + play

                                essay2 \
0 making people laugh.<br />\nranting about a go...
1 being silly. having ridiculous amonts of fun w...
2 improvising in different contexts. alternating...
3 playing synthesizers and organizing books acco...
4 creating imagery to look at:<br />\nhttp://bag...

                                essay3 ... \
0 the way i look. i am a six foot half asian, ha... ...
1                                     NaN ...
2 my large jaw and large glasses are the physica... ...
3         socially awkward but i do my best ...
4         i smile a lot and my inquisitive nature ...

                                location \
0 south san francisco, california
1         oakland, california
2         san francisco, california
3         berkeley, california
4         san francisco, california

                                offspring orientation \
0 doesn't have kids, but might want them    straight
1 doesn't have kids, but might want them    straight
2                                     NaN    straight
3         doesn't want kids    straight
4                                     NaN    straight

                                pets                                religion sex \
0 likes dogs and likes cats    agnosticism and very serious about it    m
1 likes dogs and likes cats    agnosticism but not too serious about it    m
2         has cats                                NaN    m
3         likes cats                                NaN    m
4 likes dogs and likes cats                                NaN    m

```

	sign	smokes \
0	gemini	sometimes
1	cancer	no
2	pisces but it doesn't matter	no
3	pisces	no
4	aquarius	no

	speaks	status
0	english	single
1	english (fluently), spanish (poorly), french (...)	single
2	english, french, c++	available
3	english, german (poorly)	single
4	english	single

[5 rows x 31 columns]

1.4.2 Initial Data Exploration

```
[53]: # Basic characteristics per column:
profiles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59946 entries, 0 to 59945
Data columns (total 31 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   59946 non-null  int64
1   body_type             54650 non-null  object
2   diet                  35551 non-null  object
3   drinks               56961 non-null  object
4   drugs                 45866 non-null  object
5   education             53318 non-null  object
6   essay0                54458 non-null  object
7   essay1                52374 non-null  object
8   essay2                50308 non-null  object
9   essay3                48470 non-null  object
10  essay4                49409 non-null  object
11  essay5                49096 non-null  object
12  essay6                46175 non-null  object
13  essay7                47495 non-null  object
14  essay8                40721 non-null  object
15  essay9                47343 non-null  object
16  ethnicity             54266 non-null  object
17  height               59943 non-null  float64
18  income               59946 non-null  int64
19  job                   51748 non-null  object
20  last_online          59946 non-null  object
21  location              59946 non-null  object
```

```

22  offspring      24385 non-null  object
23  orientation    59946 non-null  object
24  pets           40025 non-null  object
25  religion       39720 non-null  object
26  sex            59946 non-null  object
27  sign           48890 non-null  object
28  smokes         54434 non-null  object
29  speaks         59896 non-null  object
30  status         59946 non-null  object
dtypes: float64(1), int64(2), object(28)
memory usage: 14.2+ MB

```

The `profiles` dataset has 59946 rows and 31 columns, only 3 of them are numerical (`age`, `height` and `income`). The project will start by getting some statistics about the numerical variables, then proceeding to explore the categorical variables, and even performing some data cleaning in case necessary.

```
[54]: profiles.sign.value_counts()
```

```

[54]: gemini and it's fun to think about      1782
      scorpio and it's fun to think about      1772
      leo and it's fun to think about          1692
      libra and it's fun to think about        1649
      taurus and it's fun to think about       1640
      cancer and it's fun to think about       1597
      pisces and it's fun to think about       1592
      sagittarius and it's fun to think about  1583
      virgo and it's fun to think about        1574
      aries and it's fun to think about        1573
      aquarius and it's fun to think about     1503
      virgo but it doesn't matter             1497
      leo but it doesn't matter               1457
      cancer but it doesn't matter            1454
      gemini but it doesn't matter            1453
      taurus but it doesn't matter            1450
      libra but it doesn't matter             1408
      aquarius but it doesn't matter          1408
      capricorn and it's fun to think about   1376
      sagittarius but it doesn't matter       1375
      aries but it doesn't matter            1373
      capricorn but it doesn't matter         1319
      pisces but it doesn't matter            1300
      scorpio but it doesn't matter           1264
      leo                                     1159
      libra                                   1098
      cancer                                 1092
      virgo                                  1029
      scorpio                                1020

```

gemini	1013
taurus	1001
aries	996
pisces	992
aquarius	954
sagittarius	937
capricorn	833
scorpio and it matters a lot	78
leo and it matters a lot	66
aquarius and it matters a lot	63
cancer and it matters a lot	63
pisces and it matters a lot	62
gemini and it matters a lot	62
libra and it matters a lot	52
taurus and it matters a lot	49
sagittarius and it matters a lot	47
aries and it matters a lot	47
capricorn and it matters a lot	45
virgo and it matters a lot	41

Name: sign, dtype: int64


This column requires some cleaning since there are only 12 signs, luckily it can be noticed that the first word is always the sign. Therefore, a new column with the correct signs, called `sign_OK`, will be created.

```
[55]: # Splitting 'sign' column by space:
str_split = profiles.sign.str.split(' ')
profiles['sign_OK'] = str_split.str.get(0)
profiles.sign_OK.value_counts()
```

```
[55]: leo          4374
      gemini       4310
      libra        4207
      cancer       4206
      virgo        4141
      taurus       4140
      scorpio      4134
      aries        3989
      pisces       3946
      sagittarius  3942
      aquarius     3928
      capricorn    3573
      Name: sign_OK, dtype: int64
```

Since there are so many variables, this project believes it is worth it to analyze and plot it one by one. The first one will be `sign_OK` as it is the variable that ML model will have to predict. ####
Signs


```
[56]: # Using Seaborn bar countplot:
ax = sns.countplot(y=profiles['sign_OK'], order=profiles["sign_OK"].
                    value_counts(ascending=True).index, palette='Spectral')
ax.set(xlim=(3400,4500))
plt.show()
```



Now starting with the independent variables (the ones used as input data to predict `sign_OK`), the numerical ones will be first: `####` Age, Height, and Income

```
[57]: # Using histograms (since this is a numerical continuous variable):
plt.figure(figsize=(18,5))
ax = plt.subplot(1,3,1)
plt.hist(profiles.age, range=(0,110), bins=30, color='cadetblue',
         edgecolor="black")
plt.xlabel("Age")
plt.ylabel("Count")
ax = plt.subplot(1,3,2)
plt.hist(profiles.height, range=(0,95), bins=30, color='crimson',
         edgecolor="black")
plt.xlabel("Height")
plt.ylabel("Count")
ax = plt.subplot(1,3,3)
plt.hist(profiles.income, range=(0,1000000), bins=30, color='coral',
         edgecolor="black")
plt.xlabel("Income")
```

```
plt.ylabel("Count")
plt.show()
```




```
[58]: # Basic statistics for numerical variables:
profiles.describe()
```

```
[58]:
```

	age	height	income
count	59946.000000	59943.000000	59946.000000
mean	32.340290	68.295281	20033.222534
std	9.452779	3.994803	97346.192104
min	18.000000	1.000000	-1.000000
25%	26.000000	66.000000	-1.000000
50%	30.000000	68.000000	-1.000000
75%	37.000000	71.000000	-1.000000
max	110.000000	95.000000	1000000.000000

As seen above, there are some outliers in the data for all 3 variables. For **age** and **height**, it would be wise to eliminate **age** or **height** as 110 or 1 respectively. **income** does not need any cleaning, as there is a considerable proportion of people that have ~1,000,000 USD income. ##### Body_type


```
[59]: ax = sns.countplot(y=profiles['body_type'], order=profiles["body_type"].
    ↪value_counts(ascending=True).index, palette='Spectral')
plt.show()
# profiles.body_type.value_counts()/len(profiles)
```

It can be seen that the first 5 categories ('average', 'fit', 'athletic', 'thin', 'curvy') group ~80% of the total data.

Diet


```
[60]: ax = sns.countplot(y=profiles['diet'], order=profiles["diet"].  
      ↪value_counts(ascending=True).index, palette='Spectral')  
plt.show()  
# profiles.diet.value_counts()/len(profiles)
```



It can be seen that the first 3 categories ('mostly anything', 'anything', 'strictly anything') belong to users who are not so strict with their diets. They group ~46% of the total data.

Drinks


```
[61]: ax = sns.countplot(y=profiles['drinks'], order=profiles["drinks"].
      ↪value_counts(ascending=True).index, palette='Spectral')
plt.show()
# profiles.drinks.value_counts()/len(profiles)
```



The first category, 'socially', groups ~70% of the data. This indicates that the OKCupid users do not drink much; as 'often', 'very often', and 'desperately' group only ~10% of the data.

Drugs

```
[62]: ax = sns.countplot(x=profiles['drugs'], order=profiles["drugs"].  
      ↪value_counts(ascending=True).index, palette='Spectral')  
plt.show()  
profiles.drugs.value_counts()/len(profiles)
```




```
[62]: never          0.629300
      sometimes      0.128983
      often          0.006839
      Name: drugs, dtype: float64
```

The vast majority of the OKCupid users say that they never use drugs (~63%), yet ~25% did not answer the question.

Education

```
[63]: plt.figure(figsize=(10, 10))
      ax = sns.countplot(y=profiles['education'], order=profiles['education'].
      ↪value_counts(ascending=True).index, palette='Spectral')
      plt.show()
      # profiles.education.value_counts()/len(profiles)
```



The biggest category is ‘graduated from college/university’ which groups 40% of the data.

Ethnicity

```
[64]: profiles.ethnicity.value_counts()
```


```
[64]: white
32831
asian
6134
hispanic / latin
2823
black
2008
other
1706
...
asian, middle eastern, black, native american, hispanic / latin, white
1
asian, middle eastern, indian, hispanic / latin
1
```

```
black, native american, pacific islander, white, other
1
asian, middle eastern, black, pacific islander
1
middle eastern, black, native american, indian, pacific islander, hispanic /
latin, white, other      1
Name: ethnicity, Length: 217, dtype: int64
```

Since there are 217 different categories in the `ethnicity` variable, for the moment this variable will not be considered for the analysis.

Job

```
[65]: plt.figure(figsize=(10, 8))
      ax = sns.countplot(y=profiles['job'], order=profiles['job'].
      ↪value_counts(ascending=True).index, palette='Spectral')
      plt.show()
      # profiles.job.value_counts()/len(profiles)
```



Job-wise, the data is quite spreaded, being the biggest category ‘other’ with 13% of the data; and ‘student’ and ‘science / tech / engineering’ with 8% of the data.

Last online

```
[66]: profiles.last_online.value_counts()/len(profiles)
```

```
[66]: 2012-06-29-22-56    0.000400
      2012-06-30-22-09    0.000384
      2012-06-30-23-27    0.000384
      2012-06-30-22-56    0.000384
      2012-06-30-21-51    0.000384
      ...
      2012-05-28-12-28    0.000017
      2012-06-06-20-54    0.000017
      2012-06-30-05-50    0.000017
      2011-12-25-22-13    0.000017
      2012-06-15-23-22    0.000017
      Name: last_online, Length: 30123, dtype: float64
```

Since there are 30123 different categories in the `last_online` variable, this variable will not be considered for the analysis.

Location

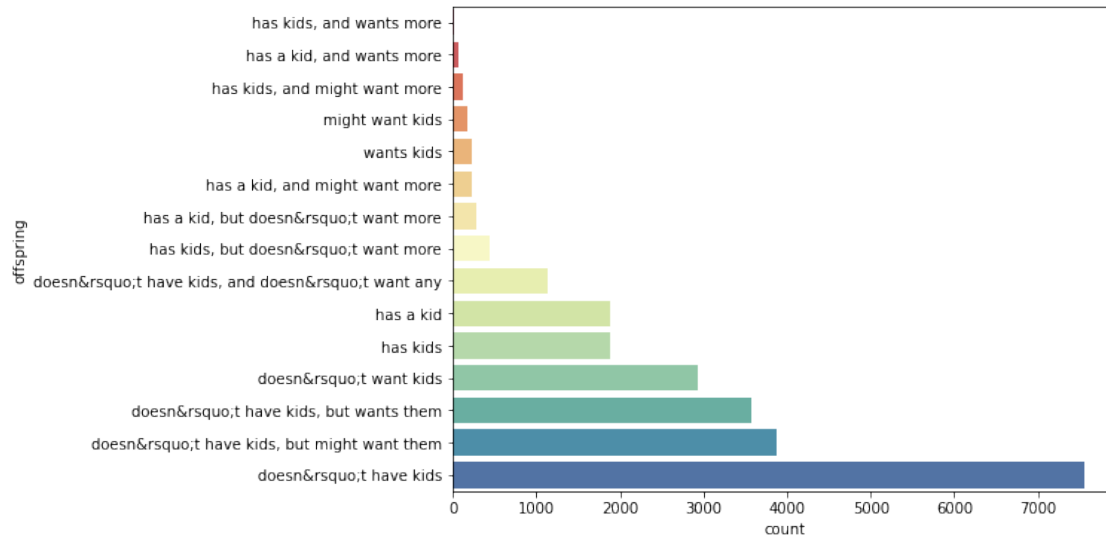
```
[67]: profiles.location.value_counts()/len(profiles)
```

```
[67]: san francisco, california    0.518200
      oakland, california         0.120342
      berkeley, california        0.070263
      san mateo, california        0.022203
      palo alto, california        0.017749
      ...
      cincinnati, ohio            0.000017
      nevada city, california      0.000017
      leander, texas               0.000017
      oakley, california           0.000017
      san luis obispo, california  0.000017
      Name: location, Length: 199, dtype: float64
```

Since there are 199 different categories in the `location` variable, for the moment we will not consider this variable for the analysis.

Offspring

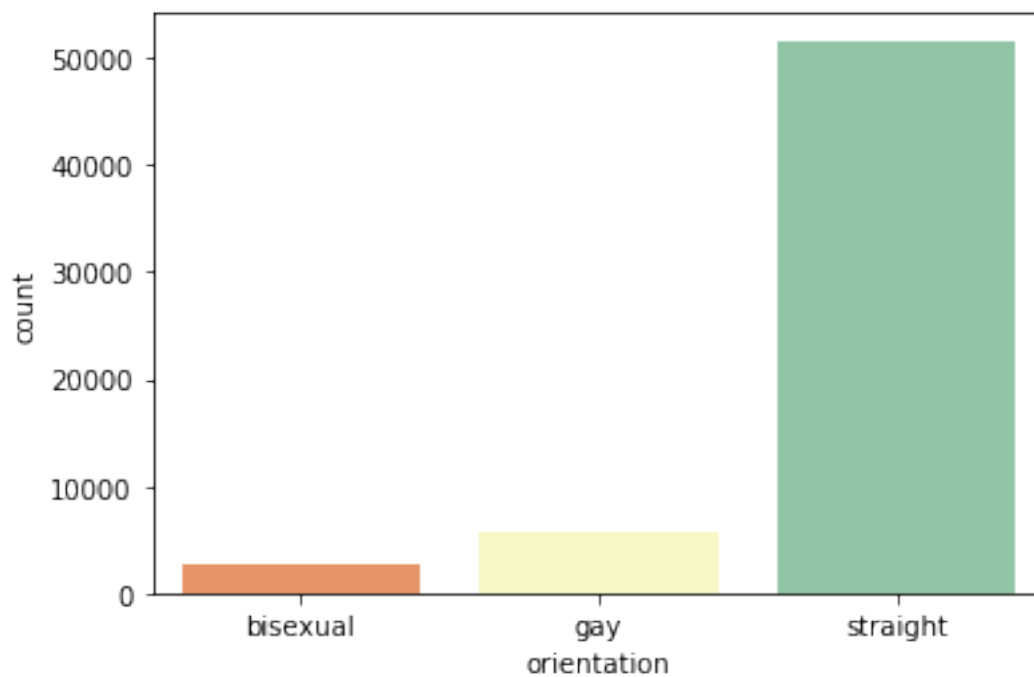
```
[68]: plt.figure(figsize=(8, 6))
      ax = sns.countplot(y=profiles['offspring'], order=profiles['offspring'].
      ↪value_counts(ascending=True).index, palette='Spectral')
      plt.show()
      # profiles.offspring.value_counts()/len(profiles)
```



The first 4 categories group ~30% of the data, corresponding to the users that do not have kids (the majority).

Orientation

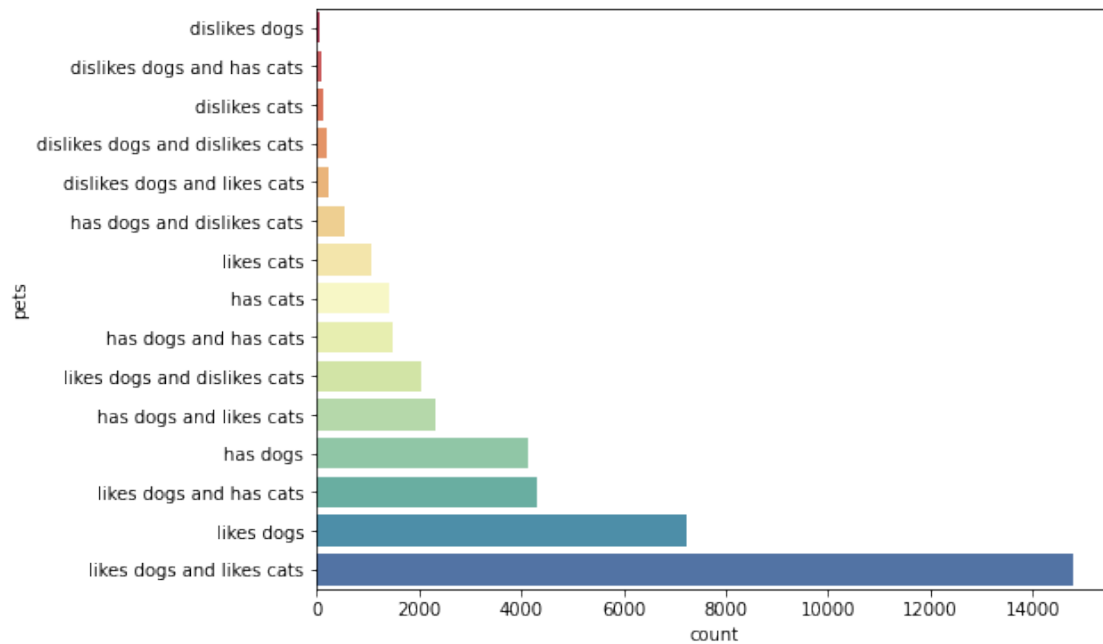
```
[69]: ax = sns.countplot(x=profiles['orientation'], order=profiles['orientation'].
      ↪ value_counts(ascending=True).index, palette='Spectral')
plt.show()
# profiles.orientation.value_counts()/len(profiles)
```



The big majority (86% of the total) have declared to be straight.

Pets


```
[70]: plt.figure(figsize=(8, 6))
      ax = sns.countplot(y=profiles['pets'], order=profiles['pets'].
      ↪value_counts(ascending=True).index, palette='Spectral')
      plt.show()
      # profiles.pets.value_counts()/len(profiles)
```



The majority of users either has dogs or likes dogs (compared to cats).


Religion

```
[71]: plt.figure(figsize=(8, 12))
      ax = sns.countplot(y=profiles['religion'], order=profiles['religion'].
      ↪value_counts(ascending=True).index, palette='Spectral')
      plt.show()
      # profiles.religion.value_counts()/len(profiles)
```



This data will be cleaned by taking only the first word, as it was done for `sign`.


```
[72]: # Splitting 'religion' column by space:
str_split_1 = profiles.religion.str.split(' ')
profiles['religion_OK'] = str_split_1.str.get(0)
# Plotting:
ax = sns.countplot(y=profiles['religion_OK'], order=profiles['religion_OK'].
    ↪value_counts(ascending=True).index, palette='Spectral')
plt.show()
# profiles.religion_OK.value_counts()/len(profiles)
```



15% of the users declare to be agnostic, and 13% do not declare any religion. Therefore, the majority does not mention affiliation with any religion. The mean age in the dataset is 32 years, which might imply some relationship with the `religion` variable, as younger generations appear less interested in it.

Sex

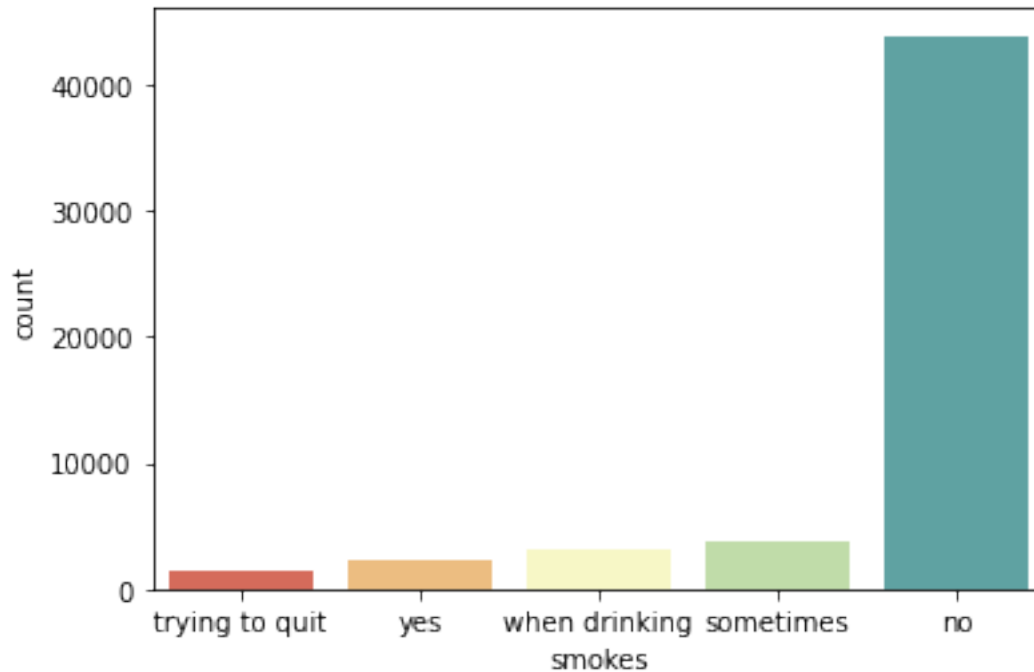
```
[73]: ax = sns.countplot(x=profiles['sex'], order=profiles['sex'].  
    ↪value_counts(ascending=True).index, palette='Spectral')  
plt.show()  
# profiles.sex.value_counts()/len(profiles)
```



60% of the users are male vs 40% who are female.

Smokes

```
[74]: ax = sns.countplot(x=profiles['smokes'], order=profiles['smokes'].  
      ↪value_counts(ascending=True).index, palette='Spectral')  
plt.show()  
# profiles.smokes.value_counts()/len(profiles)
```



More than 73% of the users declares that they do not smoke.

Speaks

```
[75]: profiles.speaks.value_counts()
```

```
[75]: english
21828
english (fluently)
6628
english (fluently), spanish (poorly)
2059
english (fluently), spanish (okay)
1917
english (fluently), spanish (fluently)
1288
...
english, spanish (okay), italian (okay), korean (poorly), hebrew (poorly)
1
english (fluently), bengali (fluently), french (fluently), spanish (okay)
1
english (fluently), spanish (poorly), tagalog (okay), french (poorly)
1
english (fluently), tagalog (okay), cebuano (okay)
1
english (fluently), spanish (okay), bengali (poorly)
```


1

Name: `speaks`, Length: 7647, dtype: int64

Since there are 7647 different categories in the `speaks` variable, this variable will not be considered for the analysis.

Status

```
[76]: ax = sns.countplot(y=profiles['status'], order=profiles['status'].  
      ↪value_counts(ascending=True).index, palette='Spectral')  
plt.show()  
# profiles.status.value_counts()/len(profiles)
```



93% of the total users are 'single', which intuitively sense since this is a dating platform.

1.4.3 Data Cleaning

Before beginning the analysis and ML classifiers design, this project will prepare/clean the data and create the required subsets. The first step is to identify the missing values in the `profiles` dataset columns (e.g. we notices that in the `drugs` column an important proportion of data was missing).

```
[77]: profiles.isnull().sum()  
np.sum(profiles.isnull())
```

```
[77]: age          0  
body_type    5296  
diet       24395
```

drinks	2985
drugs	14080
education	6628
essay0	5488
essay1	7572
essay2	9638
essay3	11476
essay4	10537
essay5	10850
essay6	13771
essay7	12451
essay8	19225
essay9	12603
ethnicity	5680
height	3
income	0
job	8198
last_online	0
location	0
offspring	35561
orientation	0
pets	19921
religion	20226
sex	0
sign	11056
smokes	5512
speaks	50
status	0
sign_OK	11056
religion_OK	20226
dtype:	int64

Now, the variables that will be considered to classify the users per zodiac sign can be chosen. The criteria for selecting them is that the variables should be related to the preferences of the users, disregarding the ones that are ‘forced’ to them by external factors or circumstances, such as **age**, **height** or **education**. A new subset that has only the selected features columns + the label column (**sign_OK**) will be created.

Note: **drugs** and **diet** were not considered due to the huge amount of blanks as per the previous information.

```
[78]: # Defining features, labels and signs list (to index answers):
features_and_labels = ['body_type', 'orientation', 'pets', 'sex', 'religion_OK',
                      'job', 'drinks', 'smokes', 'sign_OK']

signs = _
→ ['capricorn', 'aquarius', 'pisces', 'aries', 'taurus', 'gemini', 'cancer', 'leo',
   'virgo', 'libra', 'scorpio', 'sagittarius']
```

```
# Creating subset (df) with only the corresponding columns:
df = profiles[features_and_labels]
```

Since the project objective is to predict the `sign` value for the rows that do not have it, we will separate the dataframe in 2 big groups: dataset with sign (`df_with_sign`) and dataset without sign (`df_no_sign`). We will use the first dataset to train and test our models, and the second dataset will be left for trying prediction exercises if the project can find a suitable ML model.

```
[79]: # Creating the mentioned dataframes:
df_with_sign = df[df.sign_OK.notnull()]
df_no_sign = df[df.sign_OK.isnull()]
print(len(df_with_sign), len(df_no_sign))
```

```
48890 11056
```

Then, the rows from each of our subsets that do not have data in the features columns are removed, as the ML models would need this information to make the predictions.

```
[80]: # Dropping rows with empty values:
df_with_sign = df_with_sign.dropna()

# For this df we do not want to drop the empty values for the 'sign_OK' value,
↳ because the df would be empty
df_no_sign = df_no_sign.dropna(subset=[n for n in df_no_sign if n != 'sign_OK'])
print(len(df_with_sign), len(df_no_sign), '\n')

# Checking if sample used for analysis will still have representation for all
↳ zodiac signs (all OK):
print(df_with_sign.sign_OK.value_counts())
```

```
22404 2069
```

```
leo          2033
gemini       2012
virgo        1963
cancer       1920
taurus       1896
scorpio      1876
libra        1862
aries        1819
sagittarius  1794
pisces       1787
aquarius     1774
capricorn    1668
Name: sign_OK, dtype: int64
```

The selected variables are going to require some work in order to be ready for the models, mainly the creation of dummies to account for the several options available that do not follow a particular order (ordinal).

Features to convert to dummies: body_type, orientation, pets, religion_OK, sex, job, drinks, smokes.

```
[81]: # Applying One-Hot Encoding:
df_with_sign = pd.get_dummies(data=df_with_sign,
                                ↵
                                ↪columns=['body_type', 'orientation', 'pets', 'religion_OK',
                                           'sex', 'job', 'drinks', 'smokes'])
df_with_sign.head()
```

```
[81]:
```

	sign_OK	body_type_a little extra	body_type_athletic	\
0	gemini	1	0	
1	cancer	0	0	
5	taurus	0	0	
7	sagittarius	0	0	
9	cancer	0	1	

	body_type_average	body_type_curvy	body_type_fit	body_type_full figured	\
0	0	0	0	0	
1	1	0	0	0	
5	1	0	0	0	
7	1	0	0	0	
9	0	0	0	0	

	body_type_jacked	body_type_overweight	body_type_rather not say	...	\
0	0	0	0	...	
1	0	0	0	...	
5	0	0	0	...	
7	0	0	0	...	
9	0	0	0	...	

	drinks_not at all	drinks_often	drinks_rarely	drinks_socially	\
0	0	0	0	1	
1	0	1	0	0	
5	0	0	0	1	
7	0	0	0	1	
9	1	0	0	0	

	drinks_very often	smokes_no	smokes_sometimes	smokes_trying to quit	\
0	0	0	1	0	
1	0	1	0	0	
5	0	1	0	0	
7	0	1	0	0	
9	0	1	0	0	

	smokes_when drinking	smokes_yes
0	0	0

1	0	0
5	0	0
7	0	0
9	0	0

[5 rows x 74 columns]

1.4.4 Creating training and test sets

The next step is to create the training and test sets for the model from our dataset (`df_with_sign`). These sets should have 2 parts: features and labels.

```
[82]: # Defining features and labels:
features = df_with_sign[[n for n in df_with_sign if n != 'sign_OK']]
labels = df_with_sign['sign_OK']

from sklearn.model_selection import train_test_split

# Defining training and testing data and labels
train_data, test_data, train_labels, test_labels = train_test_split(features,
    ↪ labels, test_size=0.20,

    ↪ random_state=0)
print('Number of datapoints in train_data, test_data, train_labels, test_labels:
    ↪ ' +
      str(len(train_data)) + ', ' + str(len(test_data)) + ', ' +
      str(len(train_labels)) + ', ' + str(len(test_labels)))

# Transforming 'train_labels' and 'test_labels' to 1D horizontal arrays:
train_labels = train_labels.to_numpy().ravel()
test_labels = test_labels.to_numpy().ravel()
```

Number of datapoints in train_data, test_data, train_labels, test_labels: 17923, 4481, 17923, 4481

1.5 Analysis

1.5.1 K-Nearest Neighbors

The first model that the project will try is the K-Nearest Neighbors ML Classification Model. Since this model relies heavily on finding the right k-number that maximizes the accuracy score, a plot showing all scores obtained by different k-numbers (from k=1 to 40) will be created to visually select the k that maximizes the score value. Then, this k-number will be used to generate a Classification Report to assess the overall prediction score on the test data (precision, recall and overall f1-score).

```
[83]: #Creating ML K-Nearest Neighbors model:
from sklearn.neighbors import KNeighborsClassifier

k_accuracies = []
```

```


for k in range(1,41):
    k_classifier = KNeighborsClassifier(n_neighbors = k)
    k_classifier.fit(train_data, train_labels)
    k_accuracies.append(k_classifier.score(test_data, test_labels))

```

```

[84]: # Plotting accuracy vs k
k_list = range(1,41)
plt.figure(figsize=(16,5))
ax = plt.subplot()
plt.plot(range(len(k_list)), k_accuracies, linewidth=2.5)
ax.set_xticks(range(len(k_list)))
ax.set_xticklabels(k_list)
plt.xlabel('k')
plt.ylabel('Test Accuracy')
plt.title('Zodiac Sign K-Nearest Neighbors Classifier Scores')
plt.show()

```



As the graph tells, the maximum score obtained by the model is 9.5% for $k=15$ when using the model on the test data. By guessing without the help of any ML model, our accuracy would be 8.3% ($100\%/12$ possible outcomes). Therefore, it seems that the KNN model does not increase much the chance of predicting zodiac signs correctly based on the features data. Below the project shows the obtained Classification Report for $k=15$ to have more information on the precision, recall and f1-score values.

```

[85]: # Obtaining Classification Report for precision, recall and f1-score (using
      ↪ k=15):
from sklearn.metrics import classification_report
k15_classifier = KNeighborsClassifier(n_neighbors = 15).fit(train_data,
      ↪ train_labels)
k15_classifier_predictions = k15_classifier.predict(test_data)
print(classification_report(test_labels, k15_classifier_predictions))

```

```

precision    recall  f1-score   support

```

aquarius	0.10	0.17	0.12	346
aries	0.08	0.13	0.10	353
cancer	0.12	0.14	0.13	404
capricorn	0.08	0.07	0.08	334
gemini	0.11	0.11	0.11	407
leo	0.09	0.10	0.10	393
libra	0.11	0.11	0.11	351
pisces	0.08	0.07	0.08	361
sagittarius	0.10	0.07	0.09	361
scorpio	0.09	0.07	0.08	361
taurus	0.08	0.05	0.06	376
virgo	0.09	0.05	0.07	434
accuracy			0.10	4481
macro avg	0.09	0.10	0.09	4481
weighted avg	0.09	0.10	0.09	4481

The overall f1-score for the KNN model is around 10% when used for the test data. Overall this model is not recommended for predicting zodiac signs based on the given `profiles` data.

1.5.2 Decision Tree

Although the KNN Classifier model results were dissapointing, there are still other tools to try. Another option would be to use Decision Trees, which are known to be able to create accurate models for multiclass classification (which is this project's case). As seen before, one of the main challenges here is that we need to classify amongst 12 different options, so trying de Decision Tree ML model might be worth it. The tool will be tested with the same features and labels that were used for the KNN model.

Since this model relies heavily on finding the right amount of branches, the project will test via pruning (changing the `max_depth` factor) the number of branches that maximizes the accuracy score. For this, a plot chart between accuracy and `max_depth` will be created.

```
[86]: from sklearn.tree import DecisionTreeClassifier


tree_scores = []
for i in range(1,21):
    tree_classifier = DecisionTreeClassifier(random_state=1, max_depth=i)
    tree_classifier.fit(train_data, train_labels)
    score = tree_classifier.score(test_data, test_labels)
    tree_scores.append(score)

[87]: # Plotting accuracy vs k
tree_list = range(1,21)
plt.figure(figsize=(16,5))
ax = plt.subplot()
plt.plot(range(len(tree_list)), tree_scores, linewidth=2.5, color='chocolate')
ax.set_xticks(range(len(tree_list)))
```

```

ax.set_xticklabels(tree_list)
plt.xlabel('Max Depth')
plt.ylabel('Test Accuracy')
plt.title('Zodiac Sign Decision Tree Classifier Scores')
plt.show()

```



The maximum score obtained by the model is 9.1% for `max_depth=4` when using the model on the test data. Again, compared with the accuracy of random guessing (8.3%), this model also does not seem to be performing well. Below the project calculated the Classification Report for `max_depth=4` to have more information on the precision, recall and f1-score values.

```

[88]: # Obtaining Classification Report for precision, recall and f1-score (using
      ↪max_depth=4):
from sklearn.metrics import classification_report
tree_classifier = DecisionTreeClassifier(random_state=1, max_depth=4)
tree_classifier.fit(train_data, train_labels)
tree_classifier_predictions = tree_classifier.predict(test_data)
print(classification_report(test_labels, tree_classifier_predictions))

```

	precision	recall	f1-score	support
aquarius	0.00	0.00	0.00	346
aries	0.11	0.00	0.01	353
cancer	0.08	0.03	0.04	404
capricorn	0.00	0.00	0.00	334
gemini	0.09	0.20	0.13	407
leo	0.09	0.67	0.16	393
libra	0.00	0.00	0.00	351
pisces	0.00	0.00	0.00	361
sagittarius	0.07	0.02	0.03	361
scorpio	0.00	0.00	0.00	361
taurus	0.00	0.00	0.00	376
virgo	0.11	0.09	0.10	434

accuracy			0.09	4481
macro avg	0.05	0.08	0.04	4481
weighted avg	0.05	0.09	0.04	4481

```
/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

The overall f1-score for the Decision Trees model is around 9% when used for the test data. Overall this model is not recommended for predicting zodiac signs based on the given `profiles` data, in fact, it is less accurate than our previous KNN model. Still, the project will give it another go to the Decision Trees powered by Random Forests.

1.5.3 Random Forest

Despite the fact that the Decision Tree ML model score results were also dissapointing, its low accuracy could be partially related with overfitting. To discard this (and maybe come up with a better model), the project will also try to create a random forest model with 1000 trees (forest large enough to ensure that no overfitting is affecting the model) and setting features bagging to 'sqrt' as standard in ML.

```
[89]: from sklearn.ensemble import RandomForestClassifier

RF_classifier = RandomForestClassifier(n_estimators=1000, random_state=1,
↳max_features='sqrt')
RF_classifier.fit(train_data, train_labels)
RF_score = RF_classifier.score(test_data, test_labels)
print(RF_score)
```

```
0.09395224280294577
```

```
[90]: # Obtaining Classification Report for precision, recall and f1-score (using
↳k=15):
from sklearn.metrics import classification_report
RF_classifier_predictions = RF_classifier.predict(test_data)
print(classification_report(test_labels, RF_classifier_predictions))
```

	precision	recall	f1-score	support
aquarius	0.10	0.10	0.10	346
aries	0.10	0.10	0.10	353
cancer	0.10	0.11	0.11	404
capricorn	0.08	0.07	0.07	334
gemini	0.08	0.08	0.08	407
leo	0.11	0.13	0.12	393
libra	0.10	0.11	0.11	351
pisces	0.11	0.09	0.10	361
sagittarius	0.10	0.09	0.09	361
scorpio	0.07	0.06	0.06	361
taurus	0.09	0.09	0.09	376
virgo	0.09	0.08	0.09	434
accuracy			0.09	4481
macro avg	0.09	0.09	0.09	4481
weighted avg	0.09	0.09	0.09	4481

As seen above, the predictions did not get more accurate when applying Random Forest (accuracy still ~9%). This project will give it a final shot by implementing the Naive Bayes Classifier to find a better sign prediction tool than to just simply random guessing (accuracy=8.3%).

1.5.4 Naive Bayes

The Naive Bayes classifier will require a different set of data. Instead of trying to predict zodiac signs based on `body_types`, or `religion`, the project will use the essay columns to find patterns in the text written by the people in each essay to see if the NB ML model can find a better correlation between the words and the users' signs.

First step will be to prepare the data for implementing Naive Bayes.

```
[116]: # Creating function for preparing required datasets:
essay_df = profiles[['essay0', 'sign_OK']]
essay_df = essay_df.dropna()
print(len(essay_df))

# Creating dataframes (1 per sign):
aries_essays = essay_df[essay_df.sign_OK == 'aries']
taurus_essays = essay_df[essay_df.sign_OK == 'taurus']
gemini_essays = essay_df[essay_df.sign_OK == 'gemini']
cancer_essays = essay_df[essay_df.sign_OK == 'cancer']
leo_essays = essay_df[essay_df.sign_OK == 'leo']
virgo_essays = essay_df[essay_df.sign_OK == 'virgo']
libra_essays = essay_df[essay_df.sign_OK == 'libra']
scorpio_essays = essay_df[essay_df.sign_OK == 'scorpio']
sagittarius_essays = essay_df[essay_df.sign_OK == 'sagittarius']
capricorn_essays = essay_df[essay_df.sign_OK == 'capricorn']
```

```
aquarius_essays = essay_df[essay_df.sign_OK == 'aquarius']
pisces_essays = essay_df[essay_df.sign_OK == 'pisces']
```

44796

```
[117]: # Creating lists for model:
aries_text = aries_essays['essay0'].tolist()
taurus_text = taurus_essays['essay0'].tolist()
gemini_text = gemini_essays['essay0'].tolist()
cancer_text = cancer_essays['essay0'].tolist()
leo_text = leo_essays['essay0'].tolist()
virgo_text = virgo_essays['essay0'].tolist()
libra_text = libra_essays['essay0'].tolist()
scorpio_text = scorpio_essays['essay0'].tolist()
sagittarius_text = sagittarius_essays['essay0'].tolist()
capricorn_text = capricorn_essays['essay0'].tolist()
aquarius_text = aquarius_essays['essay0'].tolist()
pisces_text = pisces_essays['essay0'].tolist()

# Combining tweets into 1 list:
text_lists = []
    ↳ [aries_text,taurus_text,gemini_text,cancer_text,leo_text,virgo_text,libra_text,
        ↳
        ↳ scorpio_text,sagittarius_text,capricorn_text,aquarius_text,pisces_text]
all_texts = []
labels = []
for i in range(len(text_lists)):
    all_texts += text_lists[i]
    labels += [i]*len(text_lists[i])
print(len(all_texts))
print(len(labels))
```

44796

44796

```
[122]: # Creating the model:
train_tdata, test_tdata, train_tlabels, test_tlabels = []
    ↳ train_test_split(all_texts, labels,
        ↳
        ↳ random_state=1, test_size=0.2)
print(len(train_tdata))
print(len(test_tdata))
```

35836

8960

```
[119]: # Creating and fitting the counter:
from sklearn.feature_extraction.text import CountVectorizer
```



```

counter = CountVectorizer()
counter.fit(train_tdata)
train_tcounts = counter.transform(train_tdata)
test_tcounts = counter.transform(test_tdata)

```

```

[120]: # Creating and fitting the classifier:
from sklearn.naive_bayes import MultinomialNB
NB_classifier = MultinomialNB()
NB_classifier.fit(train_tcounts, train_tlabels)
NB_predictions = NB_classifier.predict(test_tcounts)

```

```

[121]: # Testing the model accuracy:
from sklearn.metrics import accuracy_score
print(accuracy_score(test_tlabels, NB_predictions))

```

0.09174107142857142

Again, when using the `essay0` column (essay column with the largest amount of data), the project still could not find a model with an accuracy for predicting signs beyond 9%. Regardless of this, some conclusions can be drawn and will be presented in the next section.

1.6 Evaluation

1.6.1 Conclusions

The goal of this project was to accurately predict the zodiac signs of the OKCupid users by implementing a supervised machine learning classification model. 4 different algorithms were tested (K-Nearest Neighbor, Decision Trees, Random Forests and Naive Bayes classifier), but none of them provided an acceptable improvement to just random guessing the sign without any model.

The model that achieved the highest accuracy on the test dataset was the KNN model (acc=10%). Since this results are not very satisfying, some recommendations to further investigate are given in the next section.

1.6.2 Recommendations and next steps

The main challenge posed by this project was the requirement of accurately classifying the outcomes in 12 different options (zodiac signs). For this purpose, the provided data did not prove to be enough. The recommendations for increasing the change of finding better predictions models are:

- * Increase the provided dataset to provide more data to the model to learn the differences between the features associated with 12 classes
- * Reduce the number of classes. It is not possible to reduce the number of zodiac signs, but changing the question to a yes/no binary question (like: is this user a 'libra' or not?) could make things easier for the ML models.

1.6.3 ML Models not used (and why)

- Logistic Regression: this ML tool works only for classifying between 2 possible outcomes (binary). It is possible to implement (although quite complex) by transforming the problem into a yes/no question for every sign and then comparing the probability that each datapoint has for each sign and choosing the highest one as the model prediction

- Support Vector Machines: similar to Logistic Regression, this is a powerful ML classification tool, but it works only for classifying between 2 classes (binary) and this project needed to classify amongst 12 classes